

Configurable Stack & Runtime Specialization — Addendum v0.1

Last updated: 2025-08-22

This addendum extends **Autonomous Agentic SDLC System — Master Plan (v0.1)** to make the stack fully configurable and to allow runtime specialization/assignment of agents per request.

1) Configuration Artifacts

1.1 STACK_CONFIG.yaml

```
languages: [python, typescript]
frameworks:
  web: [fastapi, express, nextjs]
  data: [pandas, polars]
db: [postgres, sqlite]
queue: [redis, rabbitmq]
deploy: [docker, k8s, ecs]
ci: [github_actions]
observability: [otel, prometheus]
security: [semgrep, trivy, gitleaks]
constraints:
  budget_usd: 50
  latency_ms_target: 500
  regions: [us-east-1]
```

1.2 TEAM_PROFILE.yaml

```
name: "Lean Web"
description: "Fast iteration for CRUD web apps"
agents:
  pm: { model: gpt-5, tools: [kb, tracker] }
  architect: { model: gpt-5, tools: [mermaid, openapi] }
  dev_backend: { model: gpt-5, tools: [pytest, uvicorn], count: 1 }
  dev_frontend: { model: gpt-5, tools: [node, playwright], count: 1 }
  qa: { model: gpt-5, tools: [pytest, coverage, playwright] }
  secops: { model: gpt-5, tools: [semgrep, trivy, gitleaks] }
  devops: { model: gpt-5, tools: [terraform, kubectl, gha] }
policies:
  approvals: { prod_deploy: "auto" }
```

```
coverage_gate: 0.8
risk_threshold: medium
```

2) Orchestration Changes

- **Agent Registry**: a catalog of agent types, their tools, costs, and rate limits.
- **Capability Graph**: skills → tools → policies; tasks are matched to agents at runtime.
- **Profile Loader**: loads selected `TEAM_PROFILE.yaml` and binds tools/models.
- **Dynamic Assignment**: planner assigns agents per task; can override by policy (e.g., high risk → enable HITL critic).
- **Budget/Latency Controller**: enforces `constraints` and triggers graceful degradation or tool swaps.

3) Runtime Selection Flow

1. User creates a **Workspace** and picks a **Profile** or uploads `STACK_CONFIG.yaml`.
2. Orchestrator validates config against support matrix; scaffolds project accordingly.
3. Tasks enter the DAG; the **Scheduler** picks agents based on capability/constraints.
4. Policies (coverage, approvals, security) are enforced as CI/CD and OPA gates.

4) Repo Additions

```
repo/
  configs/
    STACK_CONFIG.yaml          # per-workspace
    TEAM_PROFILE.yaml          # selected team profile
  orchestrator/
    registry.yaml              # agent registry
    capability_graph.yaml       # task-skills→tools mapping
```

5) APIs (Draft)

- `POST /workspaces` { name, profile_name?, stack_config? }
- `POST /workspaces/{id}/requests` { text, files[] }
- `GET /workspaces/{id}/status`
- `PUT /workspaces/{id}/profile` { profile_name }
- `PUT /workspaces/{id}/stack` { stack_config }

6) Acceptance Tests (Examples)

- Given a **Lean Web** profile, the system scaffolds FastAPI + Next.js and deploys to staging.
- Switching to **Enterprise Secure** profile enables OPA checks and artifact signing; deployments require approval.
- Increasing `coverage_gate` to 0.9 blocks merges until tests meet the threshold.

7) Next Steps

- Implement config parser + JSONSchema validation for both YAMLS.
- Build the **Profile Loader** and **Capability Graph** evaluator.
- Extend CI to read gates from `TEAM_PROFILE.yaml`.
- Ship 3 default profiles: Lean Web, Data+ML, Enterprise Secure.